

KRUSHNA PADOLE

AI/ML Engineer, Pune, Maharashtra

◇ +919960161764 ◇ [LinkedIn Profile](#) ◇ [GitHub Profile](#) ◇ [Social Links](#) ◇ krushnapadole18@gmail.com

PROFESSIONAL SUMMARY

AI/ML Engineer with 2+ years delivering production-grade ML, Deep Learning, and Generative AI solutions. Expert in PyTorch, TensorFlow, LLMs, NLP, Computer Vision, Agentic AI, and MLOps. Specialized in LLM fine-tuning, RAG architectures, model optimization, and launching scalable solutions on AWS using Docker and Kubernetes.

TECHNICAL SKILLS

ML Frameworks	PyTorch, TensorFlow, scikit-learn, XGBoost, HuggingFace, LangChain, DeepSpeed, TensorRT
Agentic AI Frameworks	LangChain, LangGraph, CrewAI, n8n, AutoGen, Multi-Agent Orchestration
AI Specializations	NLP, Computer Vision, LLMs, Neural Architecture Search, Quantum ML, Multimodal AI
Cloud/MLOps	AWS SageMaker, Docker, Kubernetes, TensorRT, DeepSpeed, DARTS, Qiskit, PennyLane
Programming	Python (Expert), SQL, PySpark, Bash Scripting, REST APIs
Data Engineering	Pandas, NumPy, Apache Airflow, BigQuery, ETL Pipelines
Core Skills	LLM Fine-tuning, RLHF, Model Quantization, Prompt Engineering, Transfer Learning

PROFESSIONAL EXPERIENCE

YASH Technologies – AI Engineer (Full-Time) July 2025 – Present

- Architected enterprise AI solutions leveraging ML, Generative AI, LLMs, and Agentic AI, delivering 40% efficiency improvement through intelligent automation and predictive analytics.
- Enhanced deep learning models securing 25% latency reduction via quantization, pruning, and TensorRT inference on AWS using Kubernetes orchestration.

Pulse Analytics – AI/ML Engineer (Full-Time) Jan 2024 – June 2025

- Developed ML solutions attaining 88% accuracy using Random Forest and XGBoost through feature engineering and hyperparameter optimization.
- Launched real-time computer vision utilizing OpenCV and YOLO for crop monitoring, slashing water usage by 35% and boosting yield by 22%.

AI/ML PROJECTS

Enterprise Agentic AI - RFP Intelligence Platform – Spearheaded RFP platform leveraging Agentic AI and Gemini Pro featuring compliance verification (85% accuracy), NLP risk assessment, and AI-driven proposal creation. Operationalized on AWS Lambda through RAG architecture, accelerating processing by 60% and improving win rate 28%.

Intelligent Document Processing with RAG – Constructed enterprise RAG utilizing LangChain and GPT-4 securing 90% accuracy. Integrated FAISS semantic search via hybrid retrieval and cross-encoder re-ranking. Released on AWS ECS incorporating Redis caching, slashing response time by 45%.

AI Mental Health Coach – Fine-tuned BERT leveraging HuggingFace for emotion classification attaining 85% precision. Established FastAPI backend featuring real-time sentiment analysis and personalized recommendations. Incorporated SHAP explainability and HIPAA-compliant encryption reaching 95% alert accuracy.

Predictive Maintenance System – Engineered bidirectional LSTM model obtaining 92% precision and 89% recall for equipment failure prediction. Rolled out on AWS EC2 featuring Streamlit dashboard and automated retraining, decreasing downtime by 40% and cutting maintenance costs by 32%.

EDUCATION

Bachelor of Science in Data Science, Savitribai Phule Pune University (**CGPA: 9.3/10.0**) 2022 - Present
Higher Senior Secondary (CBSE), Kendriya Vidyalaya VRDE 2021 - 2022

CERTIFICATIONS

- Oracle Cloud Infrastructure 2025 Generative AI Professional – **Oracle University**
- IBM Data Science Professional Certificate – **Coursera** ([Credential](#))
- Applied Data Science Lab – **WorldQuant University** ([Credential](#))
- Data Analyst Career Track – **365 Data Science** ([Credential](#))